

Gabriel West—A01782601

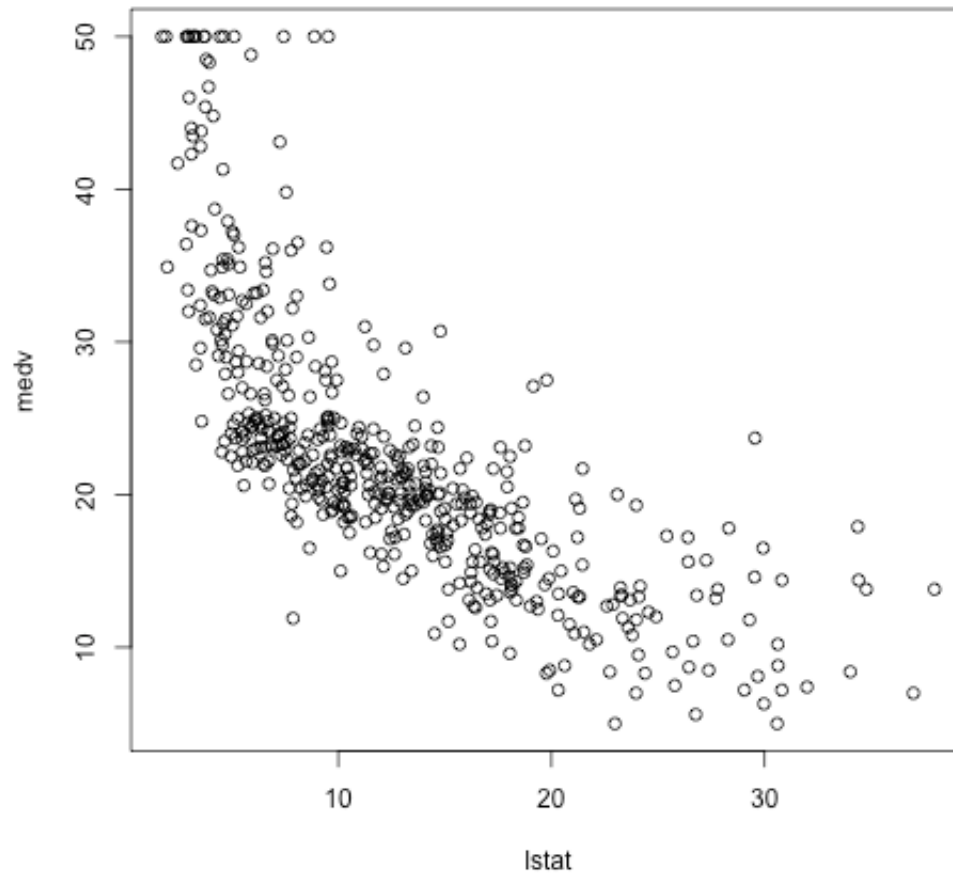
Chapter 3 Lab Thursday, March 19, 2015

In [3]: `%load_ext rmagic`

```
/usr/local/lib/python2.7/site-packages/IPython/extensions/rmagic.py:693: UserWarning: The rmagic extension in IPython is deprecated in favour of rpy2.ipython. If available, that will be loaded instead.  
http://rpy.sourceforge.net/  
    warnings.warn("The rmagic extension in IPython is deprecated in favour of "
```

In [8]: `%%R
library(MASS)
library(ISLR)
###Simple linear regression`

In [9]: `%%R
plot(medv~lstat,Boston)`



```
In [10]: %%R
fit1 <- lm(medv~lstat, data=Boston)
fit1
```

```
Call:
lm(formula = medv ~ lstat, data = Boston)
```

```
Coefficients:
(Intercept)      lstat
    34.55      -0.95
```

```
In [11]: %%R
summary(fit1)
```

```
Call:
lm(formula = medv ~ lstat, data = Boston)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-15.168	-3.990	-1.318	2.034	24.500

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.55384	0.56263	61.41	<2e-16 ***
lstat	-0.95005	0.03873	-24.53	<2e-16 ***

```
---
```

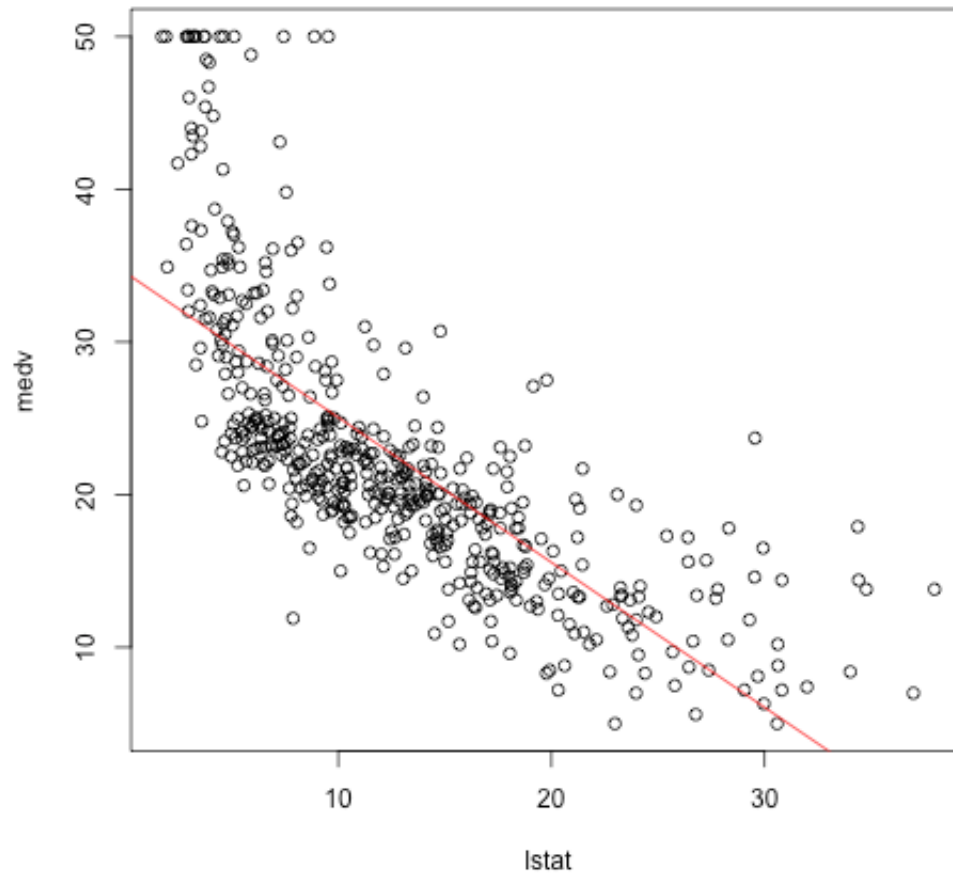
```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 6.216 on 504 degrees of freedom
```

```
Multiple R-squared:  0.5441,    Adjusted R-squared:  0.5432
```

```
F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

```
In [13]: %%R
plot(medv~lstat, Boston)
abline(fit1, col="red")
```



```
In [15]: %%R
confint(fit1)
```

```

                2.5 %      97.5 %
(Intercept) 33.448457 35.6592247
lstat      -1.026148 -0.8739505
```

```
In [16]: %%R
predict(fit1, data.frame(lstat=c(5,10,15)), interval="confidence")
```

```

      fit      lwr      upr
1 29.80359 29.00741 30.59978
2 25.05335 24.47413 25.63256
3 20.30310 19.73159 20.87461
```

```
In [18]: %%R
fit2 <- lm(medv~lstat + age, data=Boston)
summary(fit2)
```

Call:
lm(formula = medv ~ lstat + age, data = Boston)

Residuals:

	Min	1Q	Median	3Q	Max
	-15.981	-3.978	-1.283	1.968	23.158

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	33.22276	0.73085	45.458	< 2e-16	***
lstat	-1.03207	0.04819	-21.416	< 2e-16	***
age	0.03454	0.01223	2.826	0.00491	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.173 on 503 degrees of freedom
Multiple R-squared: 0.5513, Adjusted R-squared: 0.5495
F-statistic: 309 on 2 and 503 DF, p-value: < 2.2e-16

```
In [19]: %%R
fit3 <- lm(medv ~., Boston)
summary(fit3)
```

Call:

```
lm(formula = medv ~ ., data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.595	-2.730	-0.518	1.777	26.199

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.646e+01	5.103e+00	7.144	3.28e-12	***
crim	-1.080e-01	3.286e-02	-3.287	0.001087	**
zn	4.642e-02	1.373e-02	3.382	0.000778	***
indus	2.056e-02	6.150e-02	0.334	0.738288	
chas	2.687e+00	8.616e-01	3.118	0.001925	**
nox	-1.777e+01	3.820e+00	-4.651	4.25e-06	***
rm	3.810e+00	4.179e-01	9.116	< 2e-16	***
age	6.922e-04	1.321e-02	0.052	0.958229	
dis	-1.476e+00	1.995e-01	-7.398	6.01e-13	***
rad	3.060e-01	6.635e-02	4.613	5.07e-06	***
tax	-1.233e-02	3.760e-03	-3.280	0.001112	**
ptratio	-9.527e-01	1.308e-01	-7.283	1.31e-12	***
black	9.312e-03	2.686e-03	3.467	0.000573	***
lstat	-5.248e-01	5.072e-02	-10.347	< 2e-16	***

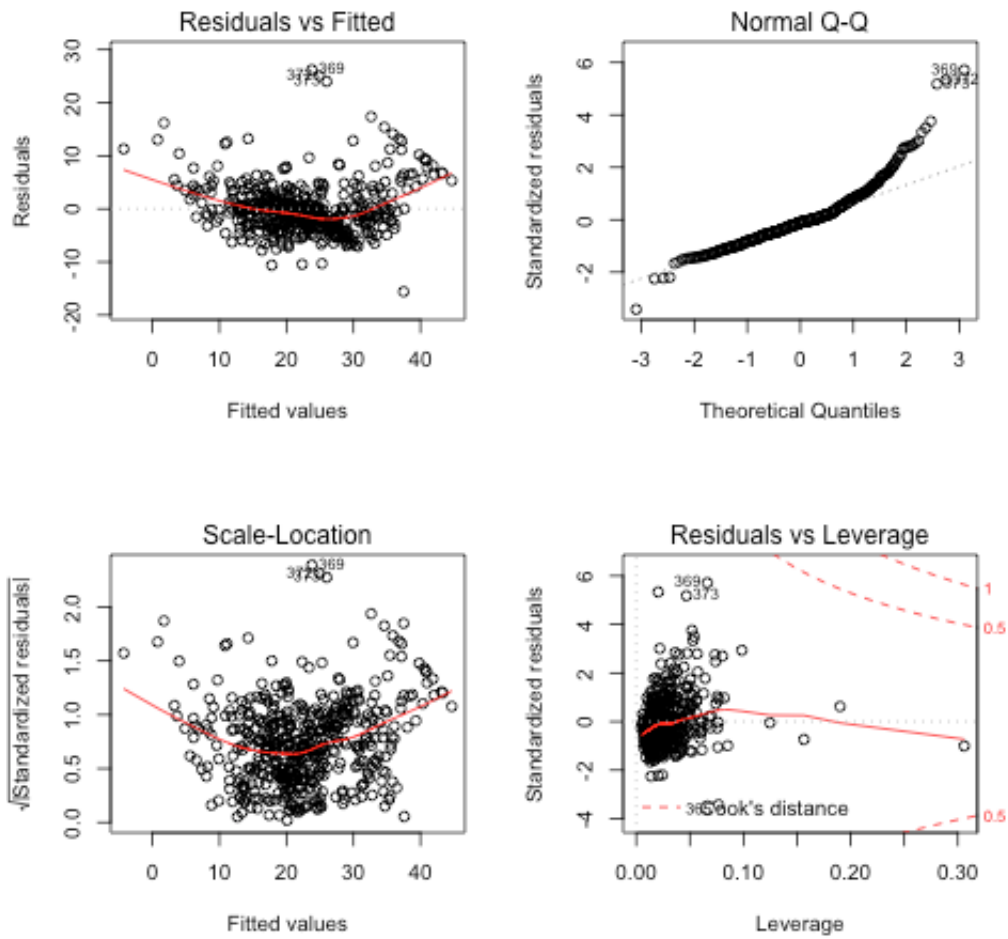
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7338

F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16

```
In [23]: %%R
par(mfrow=c(2,2))
plot(fit3)
```



```
In [25]: %%R
fit4 <- update(fit3, ~. - age - indus)
summary(fit4)
```

Call:

```
lm(formula = medv ~ crim + zn + chas + nox + rm + dis + rad +
    tax + ptratio + black + lstat, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.5984	-2.7386	-0.5046	1.7273	26.2373

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	36.341145	5.067492	7.171	2.73e-12	***
crim	-0.108413	0.032779	-3.307	0.001010	**
zn	0.045845	0.013523	3.390	0.000754	***
chas	2.718716	0.854240	3.183	0.001551	**
nox	-17.376023	3.535243	-4.915	1.21e-06	***
rm	3.801579	0.406316	9.356	< 2e-16	***
dis	-1.492711	0.185731	-8.037	6.84e-15	***
rad	0.299608	0.063402	4.726	3.00e-06	***
tax	-0.011778	0.003372	-3.493	0.000521	***
ptratio	-0.946525	0.129066	-7.334	9.24e-13	***
black	0.009291	0.002674	3.475	0.000557	***
lstat	-0.522553	0.047424	-11.019	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.736 on 494 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7348

F-statistic: 128.2 on 11 and 494 DF, p-value: < 2.2e-16


```
In [26]: %%R
fit5 <- lm(medv~lstat*age, Boston)
summary(fit5)
```

Call:

```
lm(formula = medv ~ lstat * age, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.806	-4.045	-1.333	2.085	27.552

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	36.0885359	1.4698355	24.553	< 2e-16 ***
lstat	-1.3921168	0.1674555	-8.313	8.78e-16 ***
age	-0.0007209	0.0198792	-0.036	0.9711
lstat:age	0.0041560	0.0018518	2.244	0.0252 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.149 on 502 degrees of freedom

Multiple R-squared: 0.5557, Adjusted R-squared: 0.5531

F-statistic: 209.3 on 3 and 502 DF, p-value: < 2.2e-16

```
In [27]: %%R
fit6 <- lm(medv ~ lstat + I(lstat^2), Boston)
summary(fit6)
```

Call:

```
lm(formula = medv ~ lstat + I(lstat^2), data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.2834	-3.8313	-0.5295	2.3095	25.4148

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	42.862007	0.872084	49.15	<2e-16 ***
lstat	-2.332821	0.123803	-18.84	<2e-16 ***
I(lstat^2)	0.043547	0.003745	11.63	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

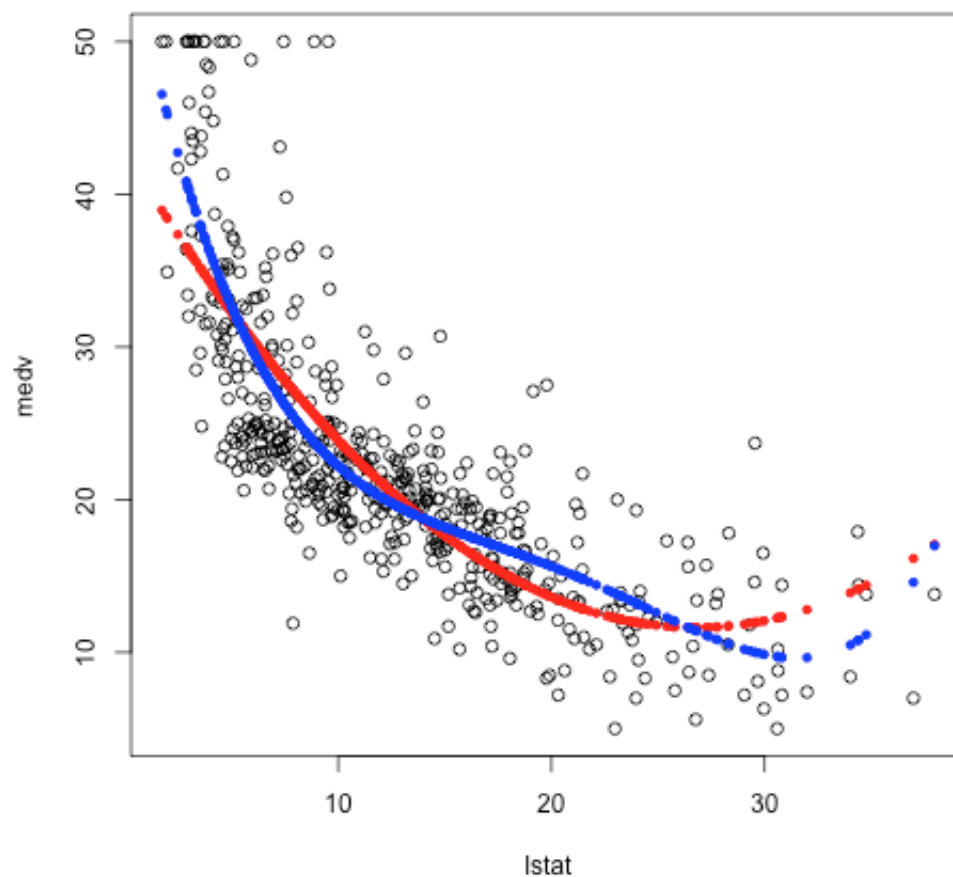
Residual standard error: 5.524 on 503 degrees of freedom

Multiple R-squared: 0.6407, Adjusted R-squared: 0.6393

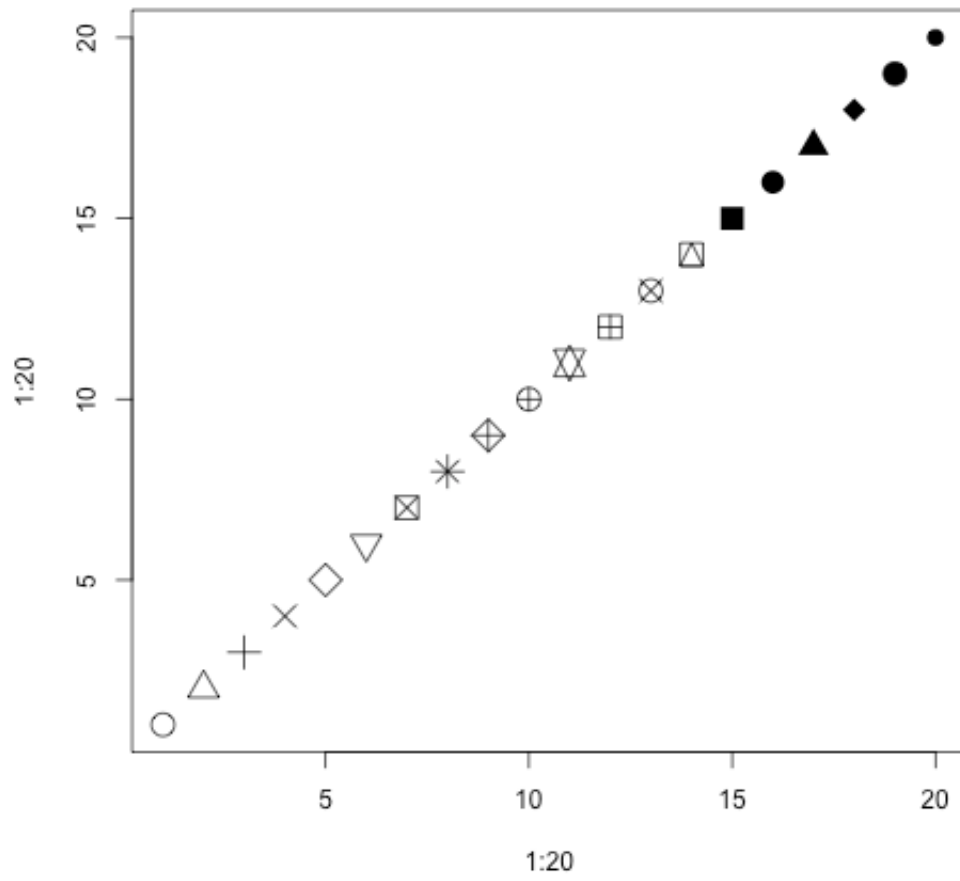
F-statistic: 448.5 on 2 and 503 DF, p-value: < 2.2e-16

```
In [33]: %%R
par(mfrow <- c(1, 1))
plot(medv ~ lstat)

points(lstat, fitted(fit6), col="red", pch=20)
fit7 <- lm(medv~poly(lstat, 4))
points(lstat, fitted(fit7), col="blue", pch=20)
```



```
In [34]: %%R
plot(1:20, 1:20, pch=1:20, cex=2)
```



```
In [37]: %%R
# Qualitative predictors
summary(Carseats)
```

Sales		CompPrice		Income		Advertising	
Min.	: 0.000	Min.	: 77	Min.	: 21.00	Min.	: 0.000
1st Qu.:	5.390	1st Qu.:	115	1st Qu.:	42.75	1st Qu.:	0.000
Median :	7.490	Median :	125	Median :	69.00	Median :	5.000
Mean :	7.496	Mean :	125	Mean :	68.66	Mean :	6.635
3rd Qu.:	9.320	3rd Qu.:	135	3rd Qu.:	91.00	3rd Qu.:	12.000
Max.	:16.270	Max.	:175	Max.	:120.00	Max.	:29.000

Population		Price		ShelveLoc		Age		Education	
Min.	: 10.0	Min.	: 24.0	Bad	: 96	Min.	:25.00	Min.	:10.0
1st Qu.:	139.0	1st Qu.:	100.0	Good	: 85	1st Qu.:	39.75	1st Qu.:	12.0
Median :	272.0	Median :	117.0	Medium:	219	Median :	54.50	Median :	14.0
Mean :	264.8	Mean :	115.8			Mean :	53.32	Mean :	13.9
3rd Qu.:	398.5	3rd Qu.:	131.0			3rd Qu.:	66.00	3rd Qu.:	16.0
Max.	:509.0	Max.	:191.0			Max.	:80.00	Max.	:18.0

Urban		US	
No	:118	No	:142
Yes:	282	Yes:	258

```
In [38]: %%R
fit1 <- lm(Sales ~. +Income:Advertising + Age:Price, Carseats)
summary(fit1)
```

Call:
lm(formula = Sales ~ . + Income:Advertising + Age:Price, data = Car
seats)

Residuals:

	Min	1Q	Median	3Q	Max
	-2.9208	-0.7503	0.0177	0.6754	3.3413

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	6.5755654	1.0087470	6.519	2.22e-10	***
CompPrice	0.0929371	0.0041183	22.567	< 2e-16	***
Income	0.0108940	0.0026044	4.183	3.57e-05	***
Advertising	0.0702462	0.0226091	3.107	0.002030	**
Population	0.0001592	0.0003679	0.433	0.665330	
Price	-0.1008064	0.0074399	-13.549	< 2e-16	***
ShelveLocGood	4.8486762	0.1528378	31.724	< 2e-16	***
ShelveLocMedium	1.9532620	0.1257682	15.531	< 2e-16	***
Age	-0.0579466	0.0159506	-3.633	0.000318	***
Education	-0.0208525	0.0196131	-1.063	0.288361	
UrbanYes	0.1401597	0.1124019	1.247	0.213171	
USYes	-0.1575571	0.1489234	-1.058	0.290729	
Income:Advertising	0.0007510	0.0002784	2.698	0.007290	**
Price:Age	0.0001068	0.0001333	0.801	0.423812	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

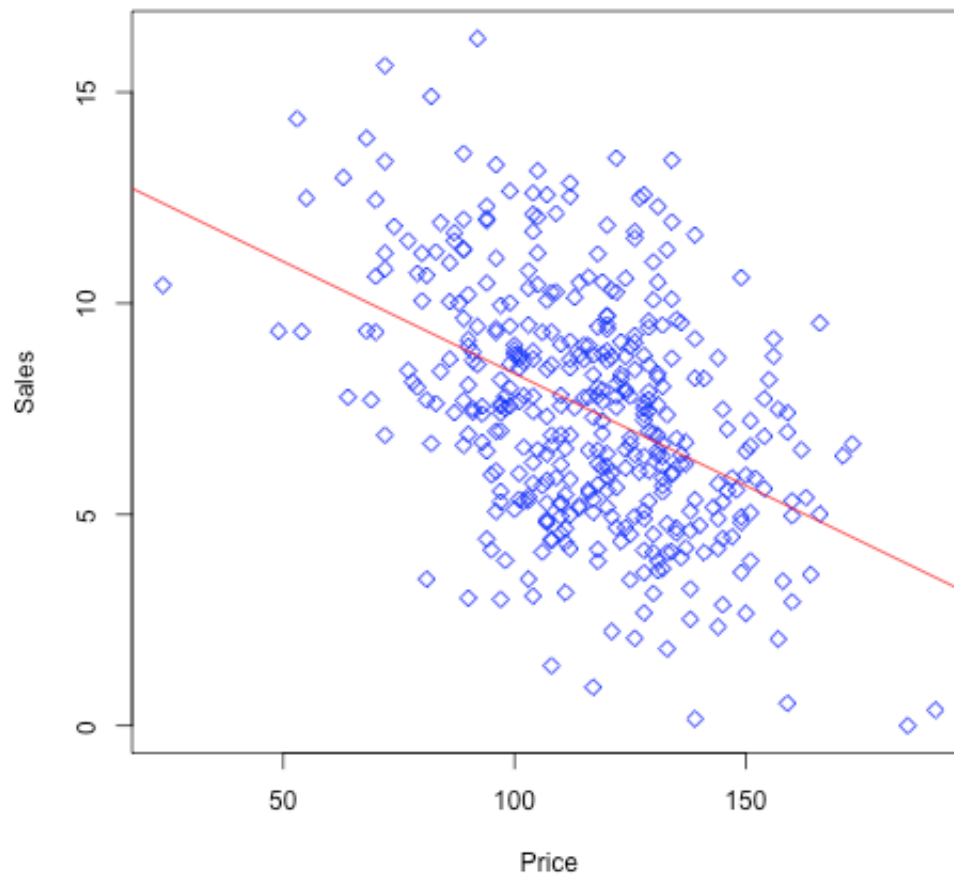
Residual standard error: 1.011 on 386 degrees of freedom
Multiple R-squared: 0.8761, Adjusted R-squared: 0.8719
F-statistic: 210 on 13 and 386 DF, p-value: < 2.2e-16

```
In [39]: %%R
contrasts(Carseats$ShelveLoc)
```

	Good	Medium
Bad	0	0
Good	1	0
Medium	0	1

```
In [48]: %%R
regplot <- function(x, y, ...){
  fit <- lm(y~x)
  plot(x, y, ...)
  abline(fit, col="red")
}

regplot(Price, Sales, xlab="Price", ylab="Sales", col="blue", pc
h=5)
```



In []: